

Apache Hadoop HBASE

Sheetal Sharma

Intern At IBM Innovation Centre

HBase is ..

- A distributed data store that can scale horizontally to 1,000s of commodity servers and petabytes of indexed storage.
- Designed to operate on top of the Hadoop distributed file system (HDFS) or Kosmos File System (KFS, aka Cloudstore) for scalability, fault tolerance, and high availability.

Benefits

- Distributed storage
- Table-like in data structure
multi-dimensional map
- High scalability
- High availability
- High performance

HBase Is Not ...

- Tables have one primary index, the *row key*.
- No join operators.
- Scans and queries can select a subset of available columns, perhaps by using a wildcard.
- There are three types of lookups:
 - ◆ Fast lookup using row key and optional timestamp.
 - ◆ Full table scan
 - ◆ Range scan from region start to end.

HBase Is Not ...(2)

- Limited atomicity and transaction support.
 - HBase supports multiple batched mutations of single rows only.
 - Data is unstructured and untyped.
- No accessed or manipulated via SQL.
 - Programmatic access via Java, REST, or Thrift APIs.
 - Scripting via JRuby.

Why HBase ?

- HBase is a Bigtable clone.
- It is open source
- It has a good community and promise for the future
- It is developed on top of and has good integration for the Hadoop platform, if you are using Hadoop already.
- It has a Cascading connector.

When to use HBase

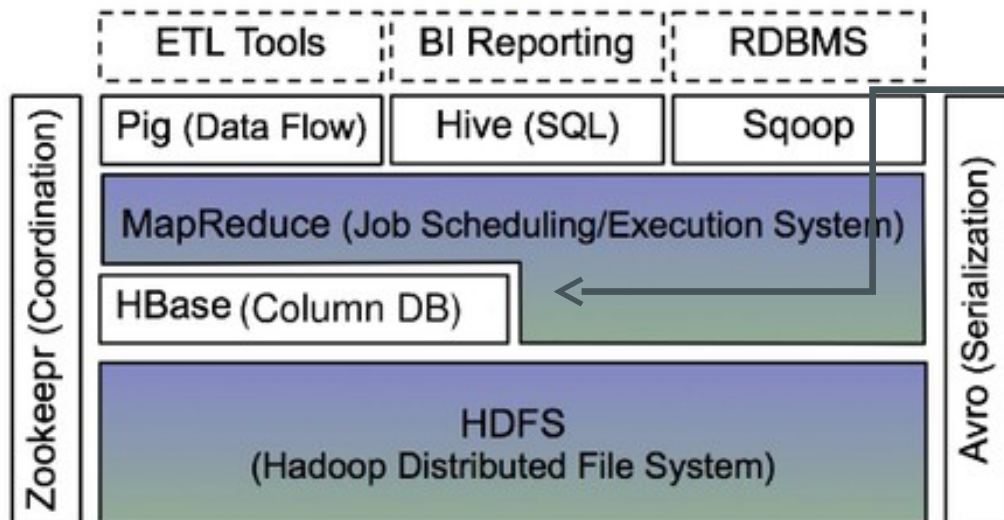
- You need random write, random read, or both (*but not neither*)
- You need to do many thousands of operations per second on multiple TB of data
- Your access patterns are well-known and simple

HBase benefits than RDBMS


- *No real indexes*
- *Automatic partitioning*
- *Scale linearly and automatically with new nodes*
- *Commodity hardware*
- *Fault tolerance*
- *Batch processing*

HBase: Part of Hadoop's Ecosystem

The Hadoop Ecosystem



HBase is built on top of HDFS



HBase files are internally stored in HDFS

HBase vs. HDFS

- Both are distributed systems that scale to hundreds or thousands of nodes
- **HDFS** is good for batch processing (scans over big files)
 - Not good for record lookup
 - Not good for incremental addition of small batches
 - Not good for updates

HBase vs. HDFS (Cont'd)

- *HBase* is designed to efficiently address the above points
 - Fast record lookup
 - Support for record-level insertion
 - Support for updates (not in place)
- HBase updates are done by creating new versions of values

HBase vs. HDFS (Cont'd)

	Plain HDFS/MR	HBase
Write pattern	Append-only	Random write, bulk incremental
Read pattern	Full table scan, partition table scan	Random read, small range scan, or table scan
Hive (SQL) performance	Very good	4-5x slower
Structured storage	Do-it-yourself / TSV / SequenceFile / Avro / ?	Sparse column-family data model
Max data size	30+ PB	~1PB

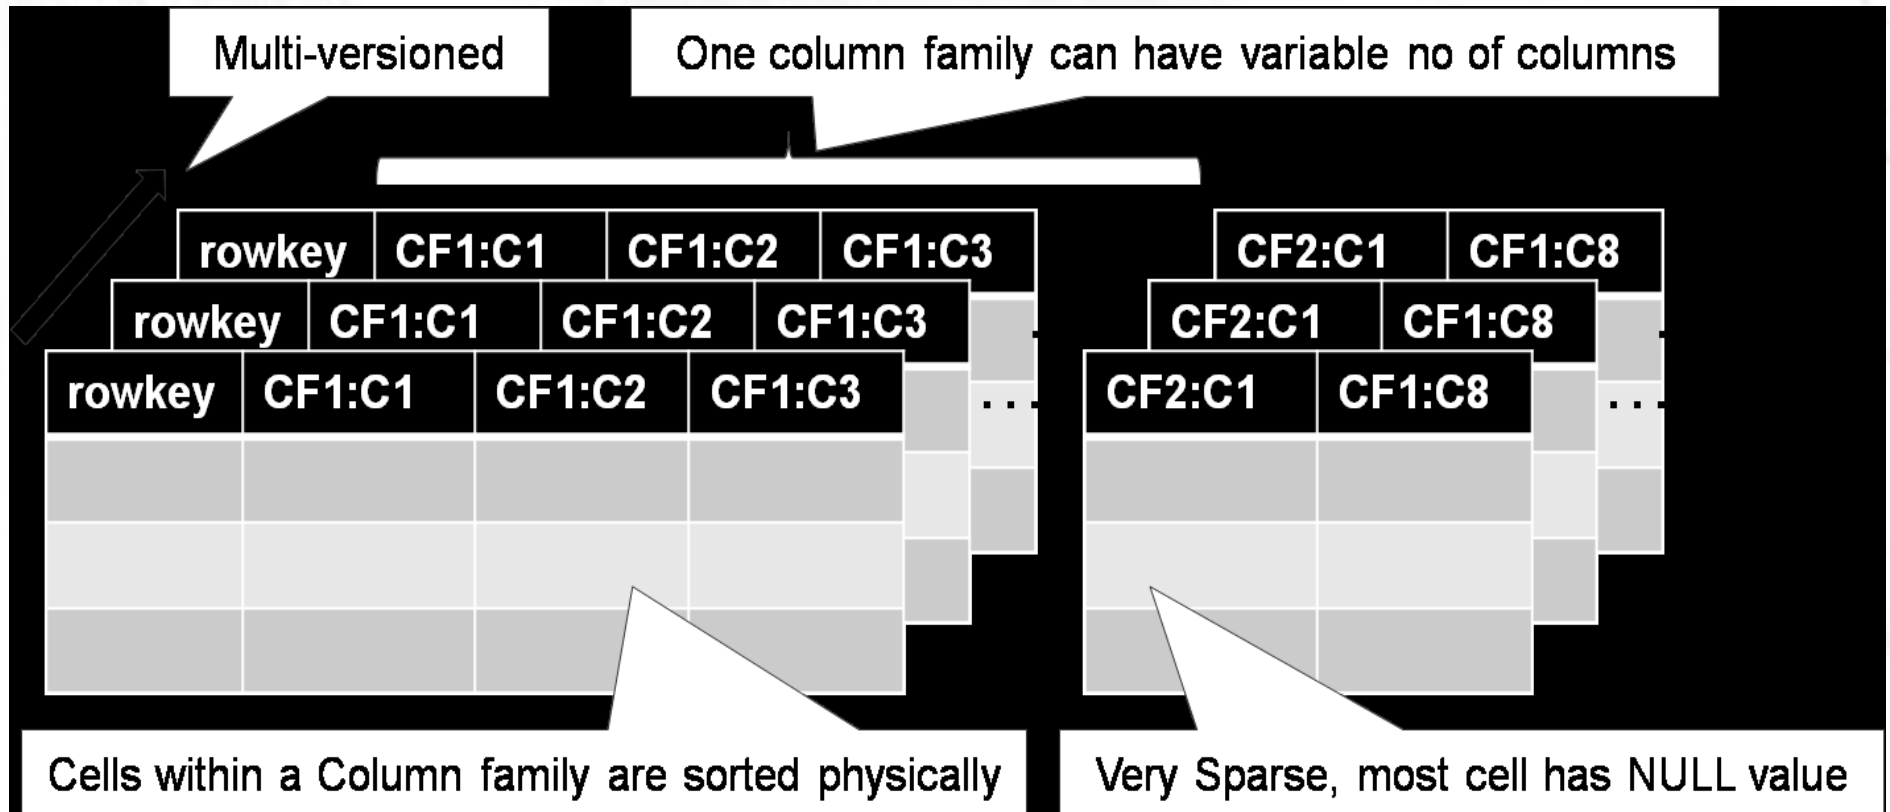
If application has neither random reads or writes ➔ Stick to HDFS

HBase vs. RDBMS

	RDBMS	HBase
Data layout	Row-oriented	Column-family-oriented
Transactions	Multi-row ACID	Single row only
Query language	SQL	get/put/scan/etc *
Security	Authentication/Authorization	Work in progress
Indexes	On arbitrary columns	Row-key only
Max data size	TBs	~1PB
Read/write throughput limits	1000s queries/second	Millions of queries/second

HBase Data Model

- Data is divided into various *tables*
- Table is composed of *columns*, columns are grouped into *column-families*



HBase Storage Model

- Partitioning
 - A table is horizontally partitioned into *regions*, each region is composed of sequential range of keys
 - Each region is managed by a *RegionServer*, a single RegionServer may hold multiple regions
- Persistence and data availability
 - HBase stores its data in HDFS, it doesn't replicate RegionServers and relies on HDFS replication for data availability.
 - Region data is cached in-memory
 - * Updates and reads are served from in-memory cache (MemStore)
 - * MemStore is flushed periodically to HDFS
 - * Write Ahead Log (stored in HDFS) is used for durability of updates

HBase: Keys and Column Families

Each record is divided into Column Families

Each row has a Key

The diagram illustrates the structure of the PERSON TABLE in HBase. It is organized into rows and columns. The first column is the 'row key', which contains 'PersonID' values ranging from 1 to 500,000,000. The subsequent columns are grouped into two main column families: 'personal_data' and 'demographic'. The 'personal_data' family contains 'Name' and 'Address' columns, while the 'demographic' family contains 'BirthDate' and 'Gender' columns. Ellipses (...) indicate that each family can contain multiple columns. Arrows from the surrounding text point to the 'row key' column and the column family headers.

PERSON TABLE					
row key	personal_data		demographic		...
PersonID	Name	Address	BirthDate	Gender	...
1	H. Houdini	Budapest, Hungary	1926-10-31	M	
2	D. Copper	New Jersey, USA	1956-09-16	M	
3	Merlin	Stonehenge, England	1136-12-03	F	
...	
500,000,000	F. Cadillac	Nevada, USA	1964-01-07	M	

Figure 2- Census Data in Column Families

Each column family consists of one or more Columns

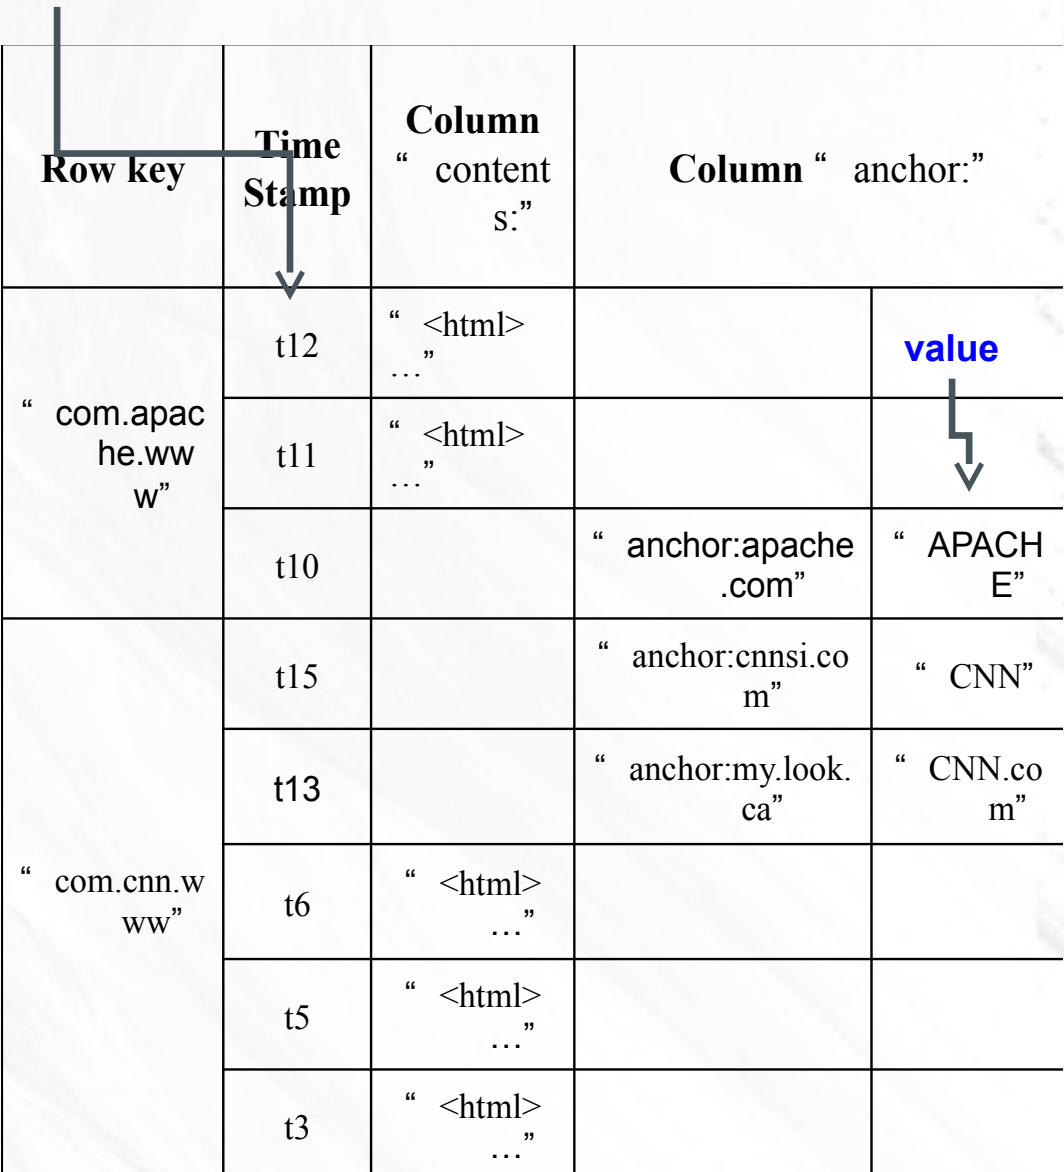
Column family named "Contents" Column family named "anchor"

- **Key**
 - Byte array
 - Serves as the primary key for the table
 - Indexed for fast lookup
- **Column Family**
 - Has a name (string)
 - Contains one or more related columns
- **Column**
 - Belongs to one column family
 - Included inside the row
 - *familyName:columnName*

Row key	Time Stamp	Column "contents:"	Column "anchor:"	
"com.apache.www"	t12	"<html>..."		
	t11	"<html>..."	Column named "apache.com"	
	t10		"anchor:apache.com"	"APACHE"
"com.cnn.www"	t15		"anchor:cnn.com"	"CNN"
	t13		"anchor:mylook.ca"	"CNN.com"
	t6	"<html>..."		
	t5	"<html>..."		
	t3	"<html>..."		

Version number for each row

- **Version Number**
 - Unique within each key
 - By default → System's timestamp
 - Data type is Long
- **Value (Cell)**
 - Byte array



Row key	Time Stamp	Column "content s:"	Column "anchor:"	
" com.apache.ww w"	t12	" <html> ... "		value
	t11	" <html> ... "		
	t10		" anchor:apache .com"	" APACHE "
" com.cnn.ww w"	t15		" anchor:cnnsi.co m"	" CNN "
	t13		" anchor:my.look. ca"	" CNN.co m "
	t6	" <html> ... "		
	t5	" <html> ... "		
	t3	" <html> ... "		

HBase Architecture

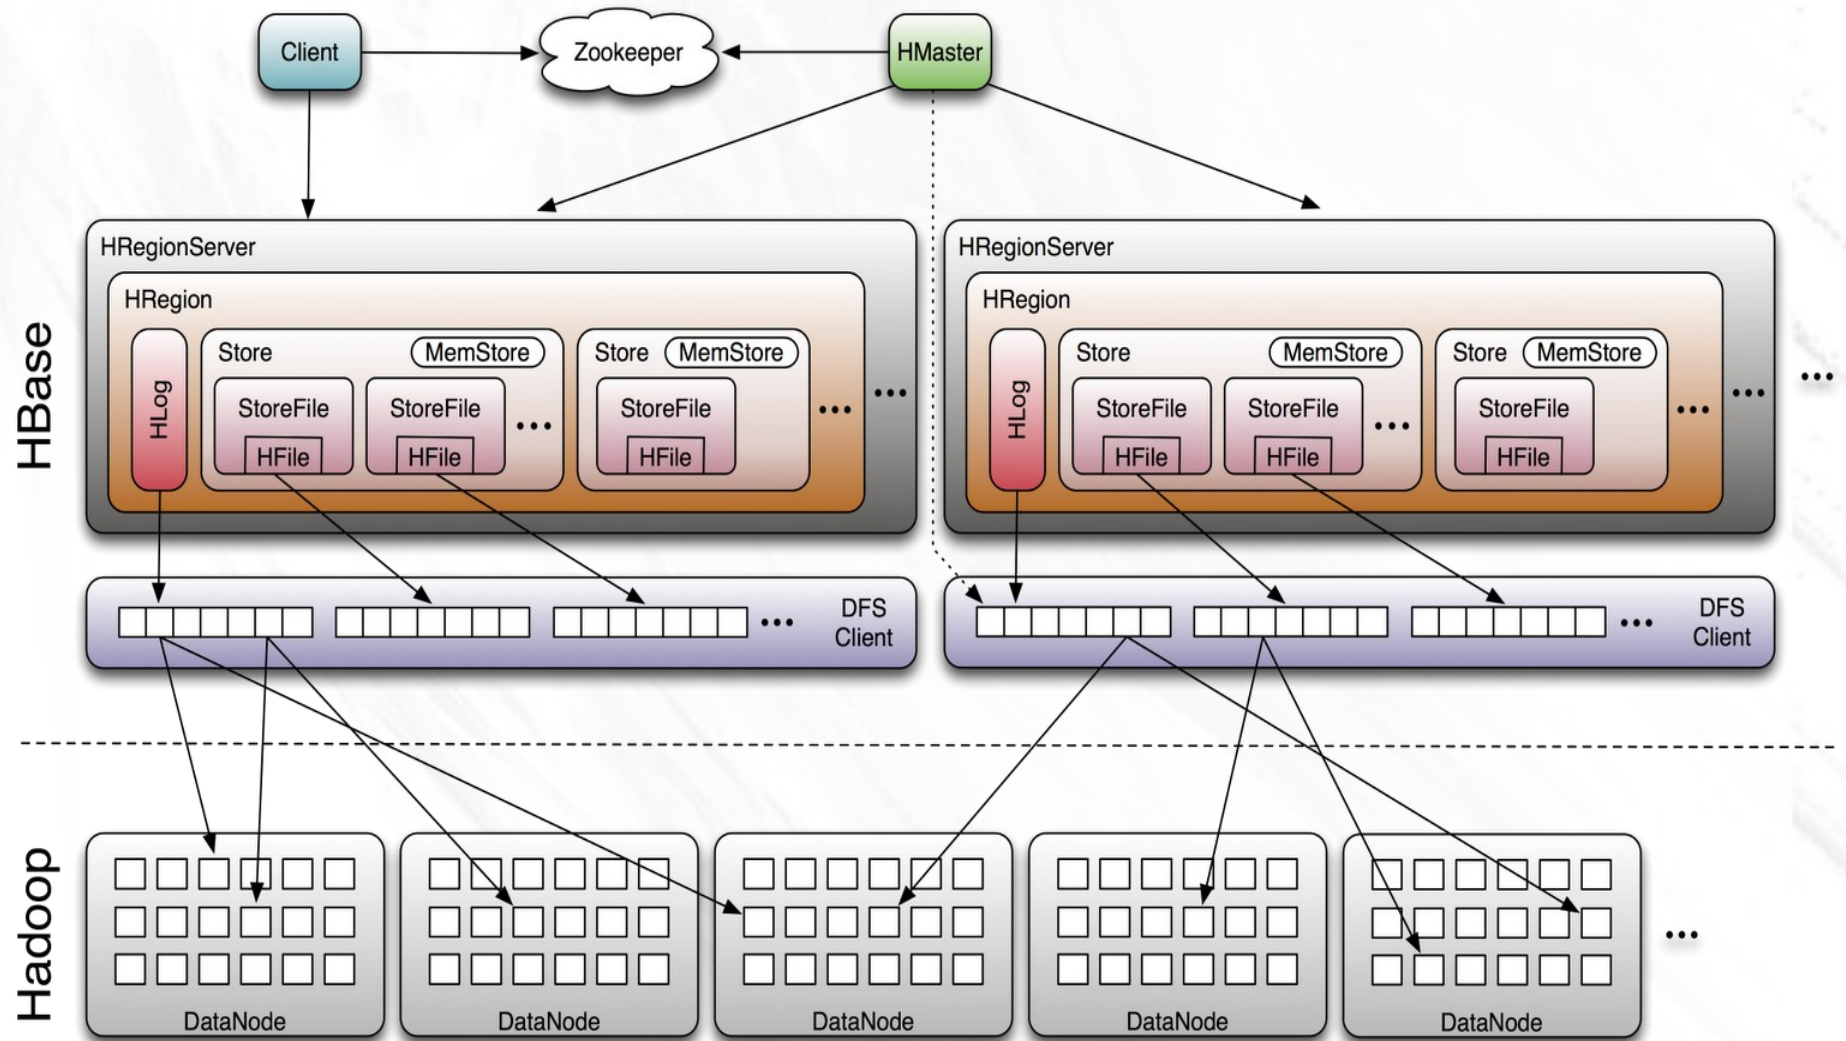
Three Major Components

- The HBaseMaster
 - One master
- The HRegionServer
 - Many region servers
- The HBase client

HBase Components


- **Region**
 - A subset of a table's rows, like horizontal range partitioning
 - Automatically done
- **RegionServer (many slaves)**
 - Manages data regions
 - Serves data for reads and writes (*using a log*)
- **Master**
 - Responsible for coordinating the slaves
 - Assigns regions, detects failures
 - Admin functions

Big Picture



ZooKeeper

- HBase depends on ZooKeeper
- By default HBase manages the ZooKeeper instance
 - E.g., starts and stops ZooKeeper
- HMaster and HRegionServers register themselves with ZooKeeper



Creating a Table

```
HBaseAdmin admin= new HBaseAdmin(config);  
HColumnDescriptor []column;  
column= new HColumnDescriptor[2];  
column[0]=new HColumnDescriptor("columnFamily1:");  
column[1]=new HColumnDescriptor("columnFamily2:");  
  
HTableDescriptor desc= new  
HTableDescriptor(Bytes.toBytes("MyTable"));  
  
desc.addFamily(column[0]);  
  
desc.addFamily(column[1]);  
  
admin.createTable(desc);
```

Operations On Regions: Get()

- Given a key → return corresponding record
- For each value return the highest version

```
Get get = new Get(Bytes.toBytes("row1"));
```

```
Result r = htable.get(get);
```

```
5.8.1.2. Default Get Example r = htable.get(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns current version of value
```

- Can control the number of versions you want

```
Get get = new Get(Bytes.toBytes("row1"));
```

```
get.setMaxVersions(3); // will return last 3 versions of row
```

```
Result r = htable.get(get);
```

```
byte[] b = r.getValue(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns current version of value
```

```
List<KeyValue> kv = r.getColumn(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns all versions of
```


Get()

Select value from table where
key='com.apache.www' AND
label='anchor:apache.com'

Row key	Time Stamp	Column "anchor:"	
"com.apache.www"	t12		
	t11		
	t10	"anchor:apache.com"	"APACHE"
	t9	"anchor:cnn.com"	"CNN"
"com.cnn.www"	t8	"anchor:my.look.ca"	"CNN.com"
	t6		
	t5		
	t3		

Operations On Regions: Scan()

```
HTable htable = ...      // instantiate HTable

Scan scan = new Scan();
scan.addColumn(Bytes.toBytes("cf"), Bytes.toBytes("attr"));
scan.setStartRow( Bytes.toBytes("row"));                // start key is inclusive
scan.setStopRow( Bytes.toBytes("row" + (char)0));      // stop key is exclusive
ResultScanner rs = htable.getScanner(scan);
try {
    for (Result r = rs.next(); r != null; r = rs.next()) {
        // process result...
    } finally {
        rs.close(); // always close the ResultScanner!
    }
}
```

Scan()

Select value from table
where anchor='cnnsi.com'

Row key	Time Stamp	Column "anchor:"	
"com.apache.www"	t12		
	t11		
	t10	"anchor:apache.com"	"APACHE"
"com.cnn.www"	t9	"anchor:cnnsi.com"	"CNN"
	t8	"anchor:my.look.ca"	"CNN.com"
	t6		
	t5		
	t3		

Operations On Regions: Put()

- Insert a new record (with a new key), Or
- Insert a record for an existing key

**Implicit version number
(timestamp)**

```
Put put = new Put(Bytes.toBytes(row));  
put.add(Bytes.toBytes("cf"), Bytes.toBytes("attr1"), Bytes.toBytes(data));  
htable.put(put);
```

Explicit version number

```
Put put = new Put(Bytes.toBytes(row));  
long explicitTimeInMs = 555; // just an example  
put.add(Bytes.toBytes("cf"), Bytes.toBytes("attr1"), explicitTimeInMs, Bytes.toBytes(data));  
htable.put(put);
```


Operations On Regions: Delete()

- Marking table cells as deleted
- **Multiple levels**
 - Can mark an entire column family as deleted
 - Can make all column families of a given row as deleted

All operations are logged by the RegionServers
The log is flushed periodically

Altering a Table

```
Configuration config = HBaseConfiguration.create();  
HBaseAdmin admin = new HBaseAdmin(conf);  
String table = "myTable";
```

```
admin.disableTable(table);
```

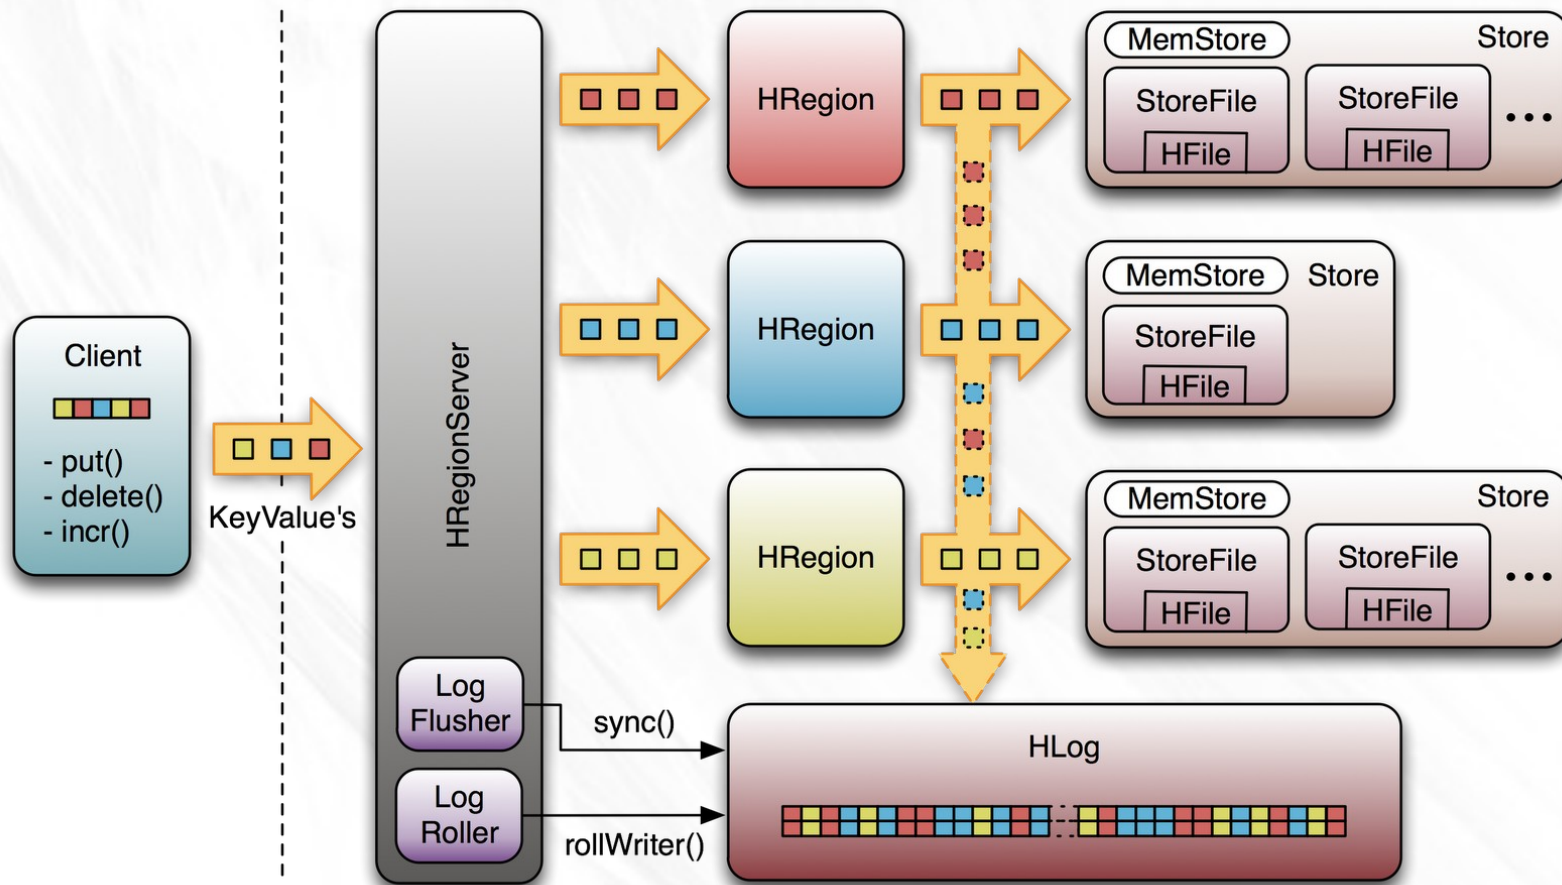
Disable the table before changing the schema

```
HColumnDescriptor cf1 = ...;  
admin.addColumn(table, cf1);    // adding new ColumnFamily  
HColumnDescriptor cf2 = ...;  
admin.modifyColumn(table, cf2); // modifying existing ColumnFamily
```

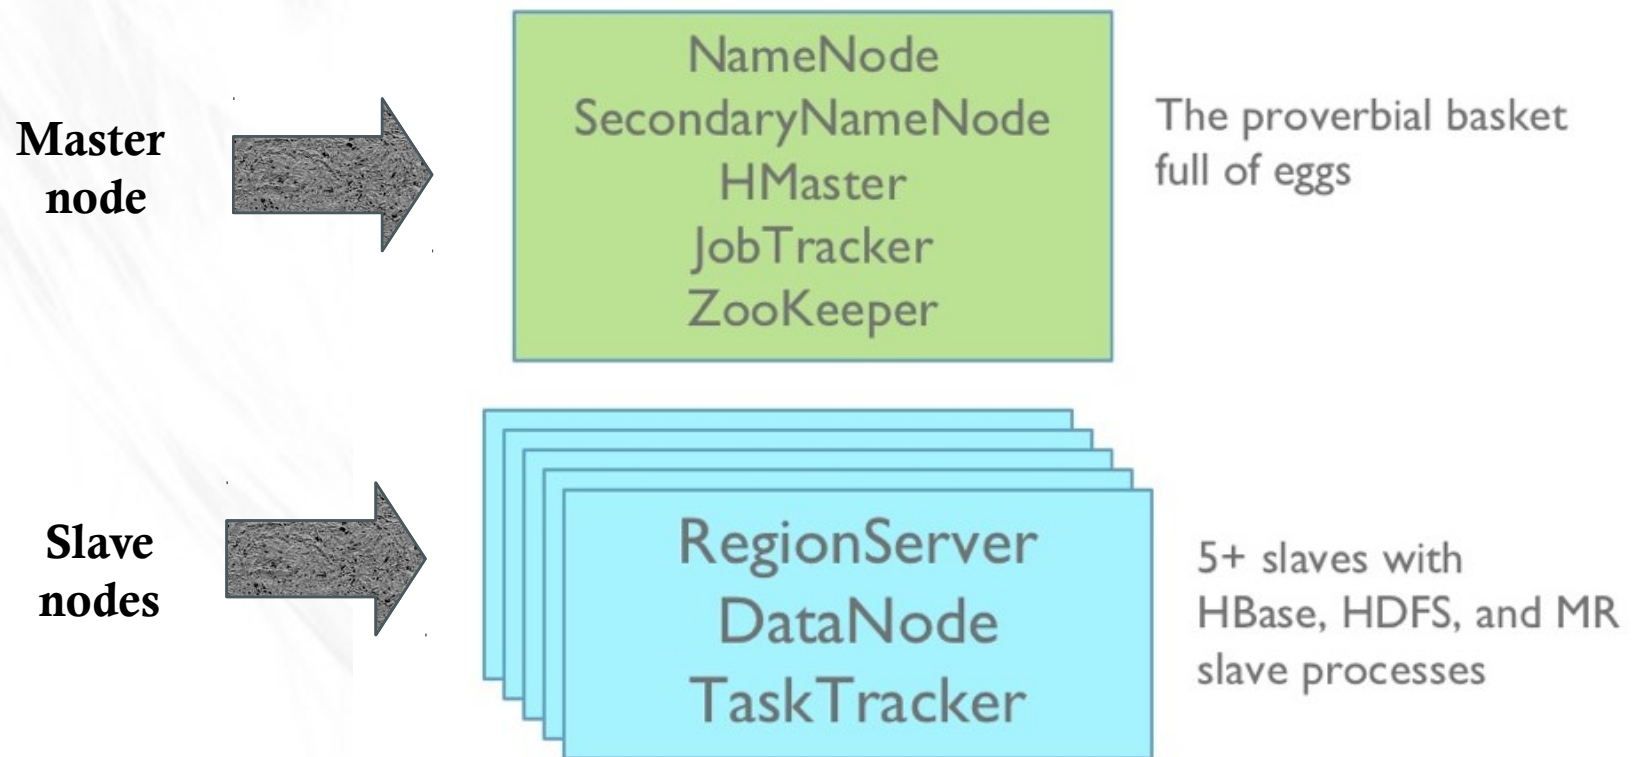
```
admin.enableTable(table);
```

6.1. Schema Creation

Logging Operations



HBase Deployment



References

- Introduction to Hbase

*trac.nchc.org.tw/cloud/raw-attachment/wiki/.../**hbase**_intro.ppt*

- *web.cs.wpi.edu/~cs525/s13-MYE/lectures/5/HBase.pptx*
- *www-users.cselabs.umn.edu/classes/Spring.../Hadoop-HBase-Tutorial.ppt*
- *www.cs.kent.edu/~jin/Cloud12Spring/HbaseHivePig.pptx*



THANK
YOU